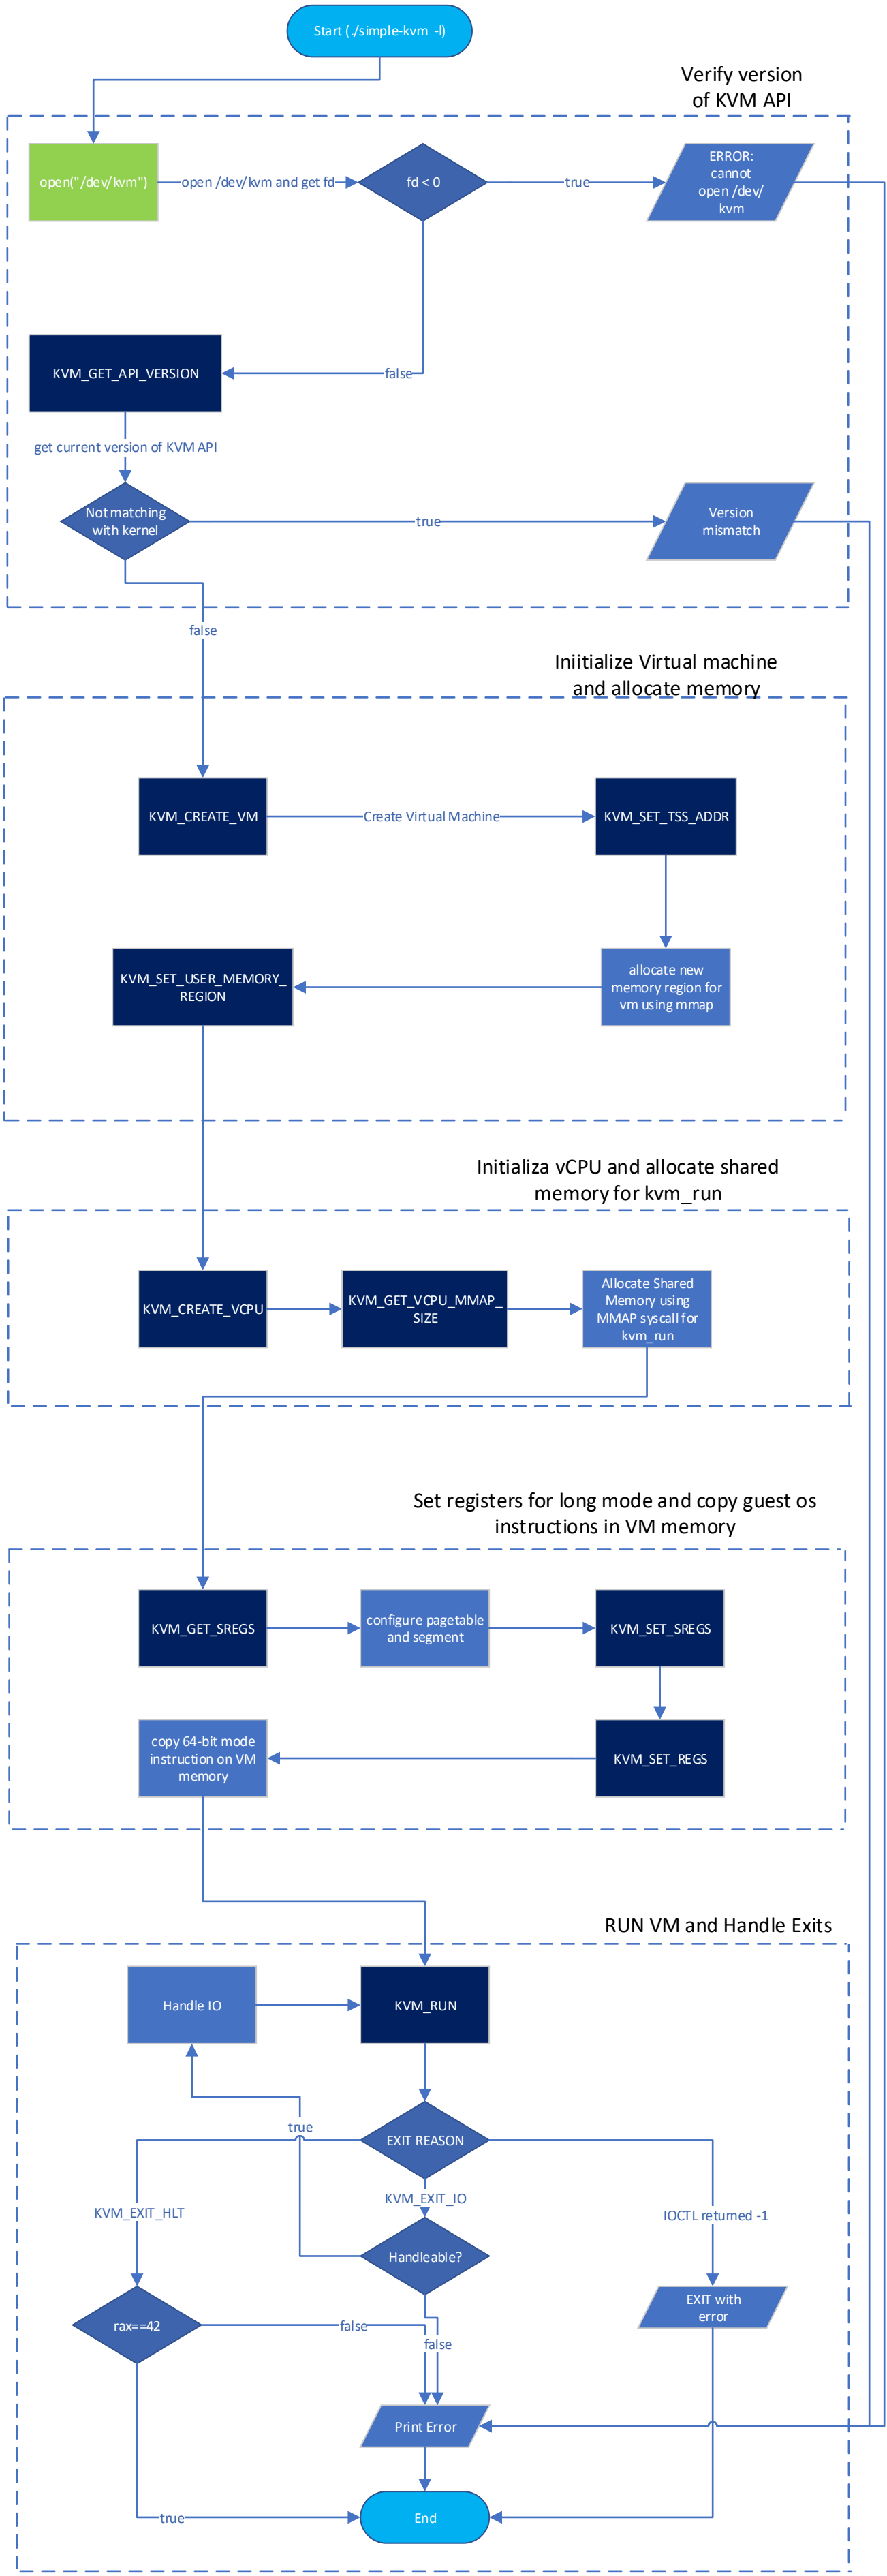


PART 1a



### 1) KVM\_GET\_API\_VERSION

- Gets KVM API version number of given system.
- IOCTL call arguments:
  - device\_fd                      Device File Descriptor(/dev/kvm)
  - KVM\_GET\_API\_VERSION          Command
  - 0                                No argument needed
- Return value:
  - integer: KVM API version number
- This IOCTL call is used to check whether underlying KVM is stable or not.

After verifying the version we want to create our own virtual machine.

### 2) KVM\_CREATE\_VM

- IOCTL call arguments:
  - device\_fd                      Device File Descriptor(/dev/kvm)
  - KVM\_CREATE\_VM                Command
  - KVM\_VM\_\*                      Machine Type e.g., KVM\_VM\_MIPS\_AUTO, From simple-kvm.c to kernel module.
- Return value:
  - integer: vm\_fd                Virtual Machine File Descriptor
- It creates an empty virtual machine, which means it doesn't have Virtual CPU and Virtual MEMORY.
- This VM can be configured using vm\_fd.

To run the above virtual machine we have to configure vCPU and vMemory.

### 3) KVM\_SET\_TSS\_ADDR

- IOCTL call arguments:
  - vm\_fd                          Virtual machine file descriptor
  - KVM\_SET\_TSS\_ADDR            Command
  - Address                       Address of virtual machine memory, from simple-kvm.c to kernel module.
- Return value:
  - integer: Zero (Success) or Non-Zero (failure)
- x86 architecture includes a special segment called "**Task State Segment**" to store hardware contexts.
- The CPU uses this address to find the kernel stack location in virtual machine memory.

After configuring TSS we will allocate memory to the virtual machine using mmap syscall. This memory block will be configured using **madvise** to scan and identify duplicate pages by KSM(Kernel Shared Memory) daemon.

### 4) KVM\_SET\_USER\_MEMORY\_REGION

- IOCTL call arguments:

- `vm_fd` Virtual machine file descriptor
  - `KVM_SET_USER_MEMORY_REGION` Command
  - `struct kvm_userspace_memory_region` Memory configuration data which includes memory address(hva), slot number, flags, starting physical address for this slot, and memory size.  
From simple-kvm.c to kernel module
- Return value:
    - integer: Zero(success) or Non-Zero(failure)
  - This IOCTL call can be used to create, modify or delete guest physical memory slots.
  - In this assignment, we are creating a guest physical memory slot.

We configured memory for VM, now we will configure vCPU for VM.

#### 5) KVM\_CREATE\_VCPU

- IOCTL call arguments:
  - `vm_fd` Virtual machine file descriptor
  - `KVM_CREATE_VCPU` Command
  - `0` No argument needed.
- Return value:
  - integer: `vcpu_fd` vCPU file descriptor which is used to configure vCPU
- This IOCTL call just creates a vCPU for a given virtual machine as we are giving `vm_fd`.
- We can only add upto **max\_vcpus** so, `vcpu_fd`  $\in$   $[0, \text{max\_vcpus})$ .

We created vCPU for a VM but now we have to assign memory for this vCPU to store and share the vCPU state with the hypervisor.

#### 6) KVM\_GET\_VCPU\_MMAP\_SIZE

- IOCTL call arguments:
  - `dev_fd` Virtual machine file descriptor
  - `KVM_GET_VCPU_MMAP_SIZE` Command
  - `0` No argument needed.
- Return value:
  - integer: `vcpu_mmap_size` Size of `struct kvm_run`, which is a shared region
- This IOCTL call just returns the size of **kvm\_run** region, this size will be used to allocate memory for that region using **mmap** syscall with **MAP\_SHARED** flag.

Mmap syscall for `kvm_run` region will be having `vcpu_fd` as an argument to let KVM know that this region belongs to which vCPU.

#### 7) KVM\_GET\_SREGS

- IOCTL call arguments:
  - `vcpu_fd` vCPU file descriptor
  - `KVM_GET_SREGS` Command
  - `struct kvm_sregs` It is an output argument from kernel module to `simple-kvm.c` which will get current special register values into pointer variable of `kvm_sregs` structure.
- Return value:
  - integer: Zero(success) or Non-Zero(failure)
- This IOCTL call gets segments registers, control registers, interrupt descriptor table address(IDT) register, global descriptor table(GDT) register, Extended feature enable register.

After getting those register values in structure pointer variable we will modify the registers to configure vCPU.

#### 8) KVM\_SET\_SREGS

- IOCTL call arguments:
  - `vcpu_fd` vCPU file descriptor
  - `KVM_SET_SREGS` Command
  - `struct kvm_sregs` Configuration for special register, which will be used by the kernel module to configure vCPU.  
From `simple-kvm.c` to kernel module.
- Return value:
  - integer: Zero(success) or Non-Zero(failure)
- This IOCTL call in `simple-kvm` just configures **control registers**, **segment registers** and **efer** as we want to run a minimal OS.

After setting special registers of vCPU, we will configure general purpose registers using below APIs.

This will be our last step of configurations, after which we'll copy compiled guest-os into **vm\_mem** area and now we are READY TO RUN!

#### 9) KVM\_GET\_REGS

- IOCTL call arguments:
  - `vcpu_fd` vCPU file descriptor
  - `KVM_GET_REGS` Command
  - `struct kvm_regs` It is an output argument from kernel module to `simple-kvm` which will get current general purpose register values into pointer variable of `kvm_regs` structure.

- Return value:
  - integer: Zero(success) or Non-Zero(failure)
- This IOCTL call gets general purpose registers in structure **kvm\_regs**.
- vCPU has 18 general purpose registers of 64 bits each.

#### 10) KVM\_SET\_REGS

- IOCTL call arguments:
 

- vcpu_fd	vCPU file descriptor
- KVM_SET_REGS	Command
- <code>struct kvm_regs</code>	Configuration for general purpose register, which will be used by the kernel module to configure vCPU. From simple-kvm.c to kernel module.
- Return value:
  - integer: Zero(success) or Non-Zero(failure)
- This IOCTL call in simple-kvm configures **rflags**, **rip** and **rsp**.

Now we are ready to launch the virtual machine using **KVM\_RUN** API.

#### 11) KVM\_RUN

- IOCTL call arguments:
 

- vcpu_fd	vCPU file descriptor
- KVM_RUN	Command
- 0	No arguments needed
- Return value:
  - integer: Zero(success) or Non-Zero(failure)
- This IOCTL call is used to actually run the VM.
- Here we are not providing any explicit argument to this IOCTL call but there is an implicit argument which is **kvm\_run**, it will be used to maintain the state of vCPU.